

**Lecture 30**

# **Inside Our Computer Vision Model**

**Mohammad Sabik Irbaz**

**Data Scientist, Leadbook Pte. Ltd.**

**Former Lead ML Engineer, Omdena & Pioneer Alpha**

**sabikirbaz@iut-dhaka.edu**

**Dokkho Data Science Career Program**

**By MasterCourse**

# *CODE*

# *Walkthrough*

```
[ ] 1 !pip install -Uqq fastai
```

```
▶ 1 from fastai.vision.all import *  
2 path = untar_data(URLs.PETS)/'images'  
3  
4 def is_cat(x): return x[0].isupper()  
5 dls = ImageDataLoaders.from_name_func(  
6     path, get_image_files(path), valid_pct=0.2, seed=42,  
7     label_func=is_cat, item_tfms=Resize(224))  
8  
9 learn = vision_learner(dls, resnet34, metrics=error_rate)  
10 learn.fine_tune(2)
```

```
▶ 1 img = PILImage.create("cat.png")  
2 img.to_thumb(192)
```

```
[ ] 1 is_cat,_,probs = learn.predict(img)  
2 print(f"Is this a cat?: {is_cat}.")  
3 print(f"Probability it's a cat: {probs[1].item():.6f}")
```

```
Is this a cat?: True.  
Probability it's a cat: 1.000000
```

```
[ ] 1 !pip install -Uqq fastai
```

```
▶ 1 from fastai.vision.all import *  
2 path = untar_data(URLs.PETS)/'images'  
3  
4 def is_cat(x): return x[0].isupper()  
5 dls = ImageDataLoaders.from_name_func(  
6     path, get_image_files(path), valid_pct=0.2, seed=42,  
7     label_func=is_cat, item_tfms=Resize(224))  
8  
9 learn = vision_learner(dls, resnet34, metrics=error_rate)  
10 learn.fine_tune(2)
```

```
▶ 1 img = PILImage.create("cat.png")  
2 img.to_thumb(192)
```

```
[ ] 1 is_cat,_,probs = learn.predict(img)  
2 print(f"Is this a cat?: {is_cat}.")  
3 print(f"Probability it's a cat: {probs[1].item():.6f}")
```

```
Is this a cat?: True.  
Probability it's a cat: 1.000000
```

Downloading and Installing the latest version of fastai.

Fastai library installs a lot of other necessary libraries like NumPy, Pandas, Torch

```
[ ] 1 !pip install -Uqq fastai
```

```
1 from fastai.vision.all import *
2 path = untar_data(URLs.PETS)/ 'images'
3
4 def is_cat(x): return x[0].isupper()
5 dls = ImageDataLoaders.from_name_func(
6     path, get_image_files(path), valid_pct=0.2, seed=42,
7     label_func=is_cat, item_tfms=Resize(224))
8
9 learn = vision_learner(dls, resnet34, metrics=error_rate)
10 learn.fine_tune(2)
```

```
1 img = PILImage.create("cat.png")
2 img.to_thumb(192)
```

```
[ ] 1 is_cat,_,probs = learn.predict(img)
2 print(f"Is this a cat?: {is_cat}.")
3 print(f"Probability it's a cat: {probs[1].item():.6f}")
```

```
Is this a cat?: True.
Probability it's a cat: 1.000000
```

Fastai has a module specifically built for vision application named "fastai.vision"

It contains different functions and libraries necessary for various computer vision models and applications.

"import \*" should be avoided when we put the code into production, but during the training and testing, it is convenient since we do not need to worry about importing necessary functions and/or libraries.

```
1 path = untar_data(URLs.PETS) / 'images'  
2 path
```

```
100.00% [811712512/811706944 01:25<00:00]  
Path('/root/.fastai/data/oxford-iiit-pet/images')
```

Download the PETs data from fastai dataset.

Extracts the data under "images" folder

docs.fast.ai/data.external.html

competition.

### Image Classification datasets

1. **CALTECH\_101**: Pictures of objects belonging to 101 categories. About 40 to 800 images per category. Most categories have about 50 images. Collected in September 2003 by Fei-Fei Li, Marco Andreetto, and Marc 'Aurelio Ranzato.
2. **CARS**: The [Cars dataset](#) contains 16,185 images of 196 classes of cars.
3. **CIFAR\_100**: The CIFAR-100 dataset consists of 60000 32x32 colour images in 100 classes, with 600 images per class.
4. **CUB\_200\_2011**: Caltech-UCSD Birds-200-2011 (CUB-200-2011) is an extended version of the CUB-200 dataset, with roughly double the number of images per class and new part location annotations
5. **FLOWERS**: 17 category [flower dataset](#) by gathering images from various websites.
6. **FOOD**:
7. **MNIST**: [MNIST dataset](#) consisting of handwritten digits
8. **PETS**: A 37 category [pet dataset](#) with roughly 200 images for each class.

```
1 len(path.ls())
```

```
7393
```

```
1 path.ls()[:5]
```

```
(#5) [Path('/root/.fastai/data/oxford-iiit-pet/images/Bombay_78.jpg'),Path('/root/.fastai/data/oxford-iiit-pet/images/chihuahua_166.jpg'),  
pet/images/havanese_73.jpg'),Path('/root/.fastai/data/oxford-iiit-pet/images/Maine_Coon_77.jpg'),Path('/root/.fastai/data/oxford-iiit-pet/
```

```
1 img = PILImage.create(path.ls()[0])  
2 img.to_thumb(192)
```



```
1 img = PILImage.create(path.ls()[1])  
2 img.to_thumb(192)
```



7393 images of different cats and dogs

```
1 len(path.ls())
```

```
7393
```

```
1 path.ls()[:5]
```

```
(#5) [Path('/root/.fastai/data/oxford-iiit-pet/images/Bombay_78.jpg'), Path('/root/.fastai/data/oxford-iiit-pet/images/chihuahua_166.jpg'),  
pet/images/havanese_73.jpg'), Path('/root/.fastai/data/oxford-iiit-pet/images/Maine_Coon_77.jpg'), Path('/root/.fastai/data/oxford-iiit-pet/
```

```
1 img = PILImage.create(path.ls()[0])  
2 img.to_thumb(192)
```



```
1 img = PILImage.create(path.ls()[1])  
2 img.to_thumb(192)
```



7393 images of different cats and dogs

Cat Image name starts with Uppercase  
Dog Image Name starts with Lowercase



Label function : returns 0 or 1

```
1 def is_cat(x): return x[0].isupper()
2 dls = ImageDataLoaders.from_name_func(
3     path,
4     get_image_files(path),
5     valid_pct=0.2,
6     seed=42,
7     label_func=is_cat,
8     item_tfms=Resize(224)
9 )
```

ImageDataLoaders.from\_name\_func [source](#)

```
ImageDataLoaders.from_name_func (path:str|Path, fnames:list,
                                label_func:callable, valid_pct=0.2,
                                seed=None, item_tfms=None,
                                batch_tfms=None, img_cls=<class
                                'fastai.vision.core.PILImage'>,
                                bs:int=64, val_bs:int=None,
                                shuffle:bool=True, device=None)
```

Create from the name attrs of `fnames` in `paths` with `label_func`

	Type	Default	Details
path	str   Path		Set the default path to a directory that a <a href="#">Learner</a> can use to save files like models
fnames	list		A list of <code>os.PathLike</code> 's to individual image files
label_func	callable		A function that receives a string (the file name) and outputs a label
valid_pct	float	0.2	
seed	NoneType	None	
item_tfms	NoneType	None	
batch_tfms	NoneType	None	
img_cls	BypassNewMeta	PILImage	
bs	int	64	Size of batch
val_bs	int	None	Size of batch for validation <a href="#">Dataloader</a>
shuffle	bool	True	Whether to shuffle data
device	NoneType	None	Device to put <a href="#">DataLoaders</a>
<b>Returns</b>	<b>DataLoaders</b>		

URL: [https://docs.fast.ai/vision.data.html#imagedataloaders.from\\_name\\_func](https://docs.fast.ai/vision.data.html#imagedataloaders.from_name_func)

```

1 def is_cat(x): return x[0].isupper()
2 dls = ImageDataLoaders.from_name_func(
3     path,
4     get_image_files(path),
5     valid_pct=0.2,
6     seed=42,
7     label_func=is_cat,
8     item_tfms=Resize(224)
9 )

```

## ImageDataLoaders.from\_name\_func [source](#)

```

ImageDataLoaders.from_name_func (path:str|Path, fnames:list,
                                label_func:callable, valid_pct=0.2,
                                seed=None, item_tfms=None,
                                batch_tfms=None, img_cls=<class
                                'fastai.vision.core.PILImage'>,
                                bs:int=64, val_bs:int=None,
                                shuffle:bool=True, device=None)

```

Create from the name attrs of `fnames` in `paths` with `label_func`

	Type	Default	Details
path	str   Path		Set the default path to a directory that a <a href="#">Learner</a> can use to save files like models
fnames	list		A list of <code>os.Pathlike</code> 's to individual image files
label_func	callable		A function that receives a string (the file name) and outputs a label
valid_pct	float	0.2	
seed	NoneType	None	
item_tfms	NoneType	None	
batch_tfms	NoneType	None	
img_cls	BypassNewMeta	PILImage	
bs	int	64	Size of batch
val_bs	int	None	Size of batch for validation <a href="#">DataLoader</a>
shuffle	bool	True	Whether to shuffle data
device	NoneType	None	Device to put <a href="#">DataLoaders</a>
<b>Returns</b>	<b>DataLoaders</b>		

```
1 learn = vision_learner(dls, resnet34, metrics=error_rate)
```

URL: [https://docs.fast.ai/vision.learner.html#vision\\_learner](https://docs.fast.ai/vision.learner.html#vision_learner)

```
1 learn = vision_learner(dls, resnet34, metrics=error_rate)
```

### vision\_learner

```
vision_learner (dls, arch, normalize=True, n_out=None, pretrained=True,
               loss_func=None, opt_func=<function Adam>, lr=0.001,
               splitter=None, cbs=None, metrics=None, path=None,
               model_dir='models', wd=None, wd_bn_bias=False,
               train_bn=True, moms=(0.95, 0.85, 0.95), cut=None,
               init=<function kaiming_normal_>, custom_head=None,
               concat_pool=True, pool=True, lin_ftrs=None, ps=0.5,
               first_bn=True, bn_final=False, lin_first=False,
               y_range=None, n_in=3)
```

Build a vision learner from `dls` and `arch`

URL: <https://fastai1.fast.ai/vision.models.html>

URL: [https://docs.fast.ai/vision.learner.html#vision\\_learner](https://docs.fast.ai/vision.learner.html#vision_learner)

### vision\_learner

```
vision_learner (dls, arch, normalize=True, n_out=None, pretrained=True,
               loss_func=None, opt_func=<function Adam>, lr=0.001,
               splitter=None, cbs=None, metrics=None, path=None,
               model_dir='models', wd=None, wd_bn_bias=False,
               train_bn=True, moms=(0.95, 0.85, 0.95), cut=None,
               init=<function kaiming_normal_>, custom_head=None,
               concat_pool=True, pool=True, lin_ftrs=None, ps=0.5,
               first_bn=True, bn_final=False, lin_first=False,
               y_range=None, n_in=3)
```

Build a vision learner from `dls` and `arch`

```
1 learn = vision_learner(dls, resnet34, metrics=error_rate)
```

### Computer Vision models zoo

The fastai library includes several pretrained models from `torchvision`, namely:

- `resnet18`, `resnet34`, `resnet50`, `resnet101`, `resnet152`
- `squeezenet1_0`, `squeezenet1_1`
- `densenet121`, `densenet169`, `densenet201`, `densenet161`
- `vgg16_bn`, `vgg19_bn`
- `alexnet`

URL: <https://fastai1.fast.ai/vision.models.html>

URL: [https://docs.fast.ai/metrics.html#error\\_rate](https://docs.fast.ai/metrics.html#error_rate)

URL: [https://docs.fast.ai/vision.learner.html#vision\\_learner](https://docs.fast.ai/vision.learner.html#vision_learner)

## vision\_learner

```
vision_learner (dls, arch, normalize=True, n_out=None, pretrained=True,
               loss_func=None, opt_func=<function Adam>, lr=0.001,
               splitter=None, cbs=None, metrics=None, path=None,
               model_dir='models', wd=None, wd_bn_bias=False,
               train_bn=True, moms=(0.95, 0.85, 0.95), cut=None,
               init=<function kaiming_normal_>, custom_head=None,
               concat_pool=True, pool=True, lin_ftrs=None, ps=0.5,
               first_bn=True, bn_final=False, lin_first=False,
               y_range=None, n_in=3)
```

Build a vision learner from `dls` and `arch`

1 learn = vision\_learner(dls, resnet34, metrics=error\_rate)

## Computer Vision models zoo

The fastai library includes several pretrained models from [torchvision](#), namely:

- resnet18, resnet34, resnet50, resnet101, resnet152
- squeezenet1\_0, squeezenet1\_1
- densenet121, densenet169, densenet201, densenet161
- vgg16\_bn, vgg19\_bn
- alexnet

## error\_rate [🔗](#)

```
error_rate (inp, targ, axis=-1)
```

1- [accuracy](#)

```
x = torch.randn(4,5)
y = x.argmax(dim=1)
test_eq(error_rate(x,y), 0)
y1 = change_targ(y, 2, 5)
test_eq(error_rate(x,y1), 0.5)
test_eq(error_rate(x.unsqueeze(1).expand(4,2,5), torch.stack([y,y1], dim=1)), 0.25)
```

```
1 learn.fine_tune(2)
```

epoch	train_loss	valid_loss	error_rate	time
-------	------------	------------	------------	------

0	0.158167	0.033317	0.009472	00:53
---	----------	----------	----------	-------

epoch	train_loss	valid_loss	error_rate	time
-------	------------	------------	------------	------

0	0.058523	0.038277	0.007442	00:51
---	----------	----------	----------	-------

1	0.029566	0.018829	0.006089	00:51
---	----------	----------	----------	-------

1. Adds custom layers
2. Freezes the pretrained layers
3. Train for one epoch
4. Unfreeze the pretrained layers
5. Train for "n" epochs

# Inference

```
1 learn.predict(img)


('True', TensorBase(1), TensorBase([3.8827e-12, 1.0000e+00]))
```

```
1 is_cat,_,probs = learn.predict(img)
2 print(f"Is this a cat?: {is_cat}.")
3 print(f"Probability it's a cat: {probs[1].item():.6f}")
```

```
Is this a cat?: True.
Probability it's a cat: 1.000000
```




# What Our Model Learns

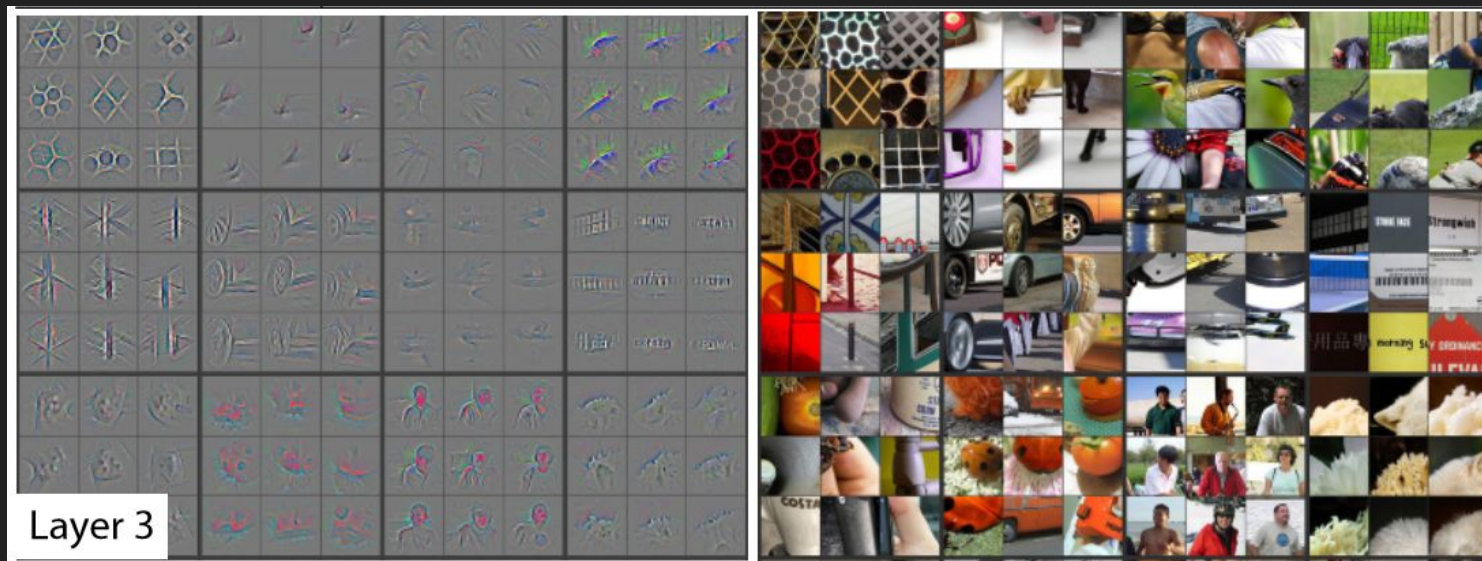


From “Visualizing and Understanding Convolutional Networks” Paper

# What Our Model Learns



# What Our Model Learns



# What Our Model Learns

